

OmniBrain Vision Test Document

Purpose: validate image/table extraction, PostgreSQL visual metadata, page selection, and Vision Agent routing.

Page 1 - Revenue Chart

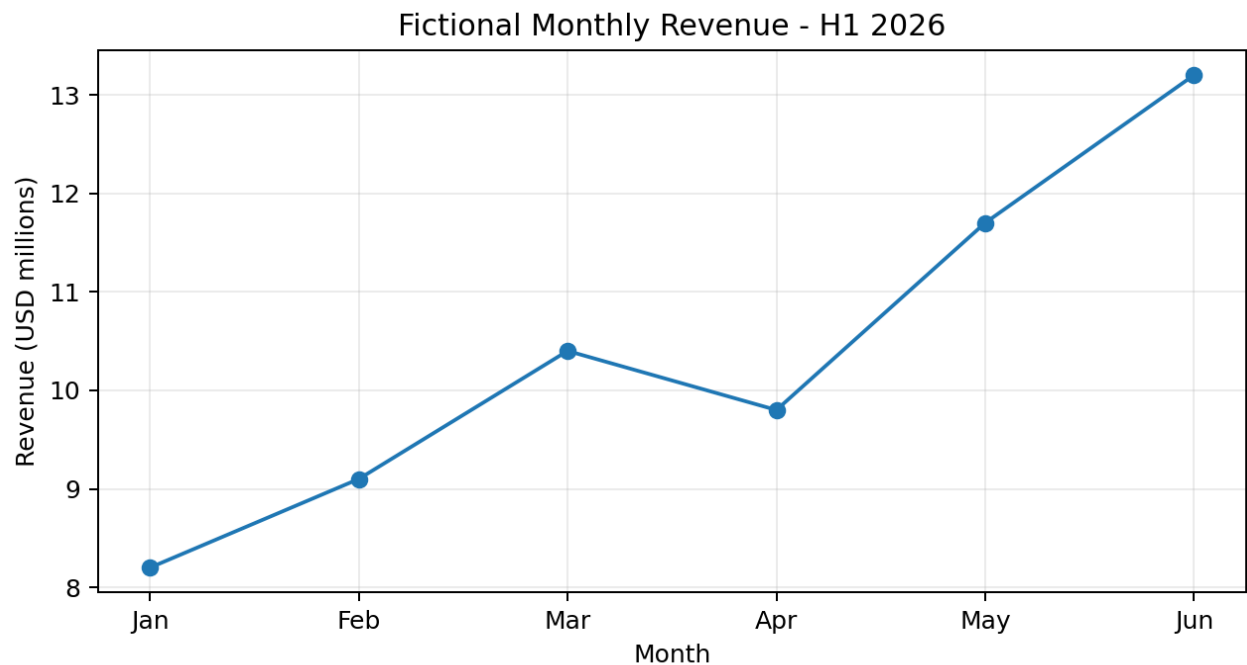


Figure 1 caption: fictional monthly revenue rises overall from USD 8.2 million in January to USD 13.2 million in June, with a temporary decline from March to April. June is the highest month.

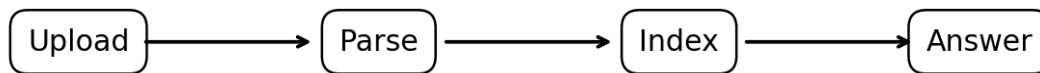
Page 2 - Regional Performance Table

Region	2025 Revenue	2026 Revenue	Growth	Customer Count
North	\$18.0M	\$21.6M	20%	410
South	\$14.5M	\$15.2M	4.8%	305
East	\$9.8M	\$12.7M	29.6%	280
West	\$11.2M	\$10.6M	-5.4%	245

Table 1 summary: East has the highest growth rate at 29.6 percent. North has the highest 2026 revenue at USD 21.6 million and the largest customer count at 410. West is the only region with negative growth, at -5.4 percent.

Page 3 - Pipeline Diagram

OmniBrain Test Pipeline



Expected flow: PDF upload -> parsing -> vector/metadata indexing -> routed answer

Figure 2 caption: the pipeline proceeds in this order: Upload, Parse, Index, Answer.

Test note: if your parser stores image captions/OCR/table summaries in PostgreSQL, the Vision Agent should be able to explain these visuals from metadata. Direct pixel-level vision requires a multimodal model and access to the stored image bytes.